

YOUR SUPPORT TEAM'S SECOND BRAIN

ML-powered answers from your own knowledge base — in under 25 ms, without sending a single query to the cloud.

Local LLM inference plus a hybrid ML search pipeline that drafts accurate, KB-grounded IT support responses entirely on the operator's laptop. SQLCipher AES-256 at rest, a self-improving feedback loop, and an ops surface for deploy, rollback, and eval — all shipped as a single Tauri app.

FIVE FEATURE PILLARS

01

ML intent classification

A calibrated logistic-regression classifier routes every ticket — policy, howto, access, incident — in under 5 ms.

0.914 macro-F1

02

Sub-25 ms hybrid search

TF-IDF retrieval plus a ms-marco-MiniLM-L-6-v2 cross-encoder reranker — citations are real files, not hallucinations.

22 ms p50 · 46 ms p95

03

Trust-gated responses

Confidence modes (answer · clarify · abstain) plus inline [n] citations link every sentence back to its source.

0.93 grounded · 0.96 fai...

04

Self-improving feedback loop

Low-confidence queries get clustered into KB gaps, ranked by impact, and turned into a prioritized list of articles.


Mock gaps · demo

05

Local-first & encrypted

SQLCipher AES-256 at rest, local Ollama inference, and a no-cloud demo path. Ships with deploy / rollback / eval ops.

SQLCipher · Ollama · Tau...



TICKET DEFLECTION

25%

of tier-1 tickets never reach a human — drafts accepted and pasted into Jira as-is.

illustrative support-automation benchmark

HYBRID SEARCH LATENCY

<25ms

p50 across TF-IDF candidate retrieval plus cross-encoder reranker — instant in the composer.

demo latency target for the local retrieval path

DEMO KB INDEXED

27

sanitized markdown policies, procedures, and references used for portfolio capture.

checked-in fake KB · no private data